

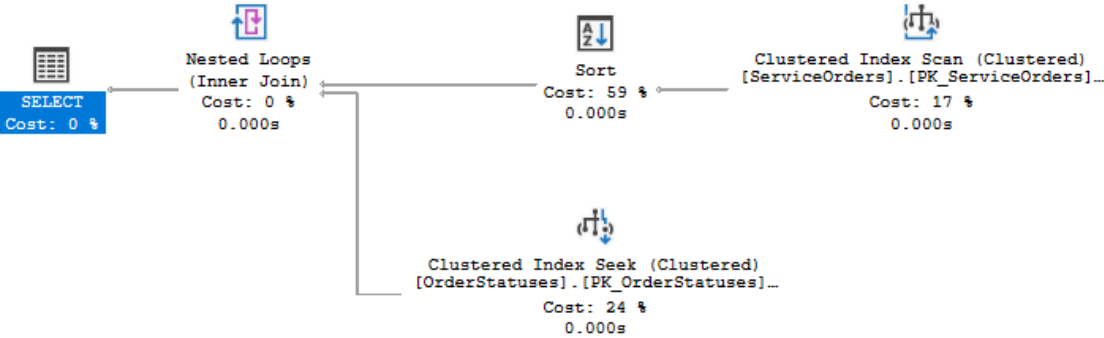
# RAPORT INDEKSY

## Zapytanie z otwartymi zleceniami serwisowymi

```
SELECT
    o.Id,
    o.OrderNumber,
    o.OpenedAt,
    o.CustomerNameSnapshot,
    o.RegistrationNumberSnapshot,
    s.Name AS StatusName
FROM dbo.ServiceOrders AS o
INNER JOIN dbo.OrderStatuses AS s
    ON s.Id = o.StatusId
WHERE s.Name <> N'Completed'
ORDER BY o.OpenedAt DESC;
```

### Bez indeksu

Query 1: Query cost (relative to the batch): 100%  
SELECT o.Id, o.OrderNumber, o.OpenedAt, o.CustomerNameSnapshot, o.RegistrationNumberSnapshot, s.Name AS Sta



Physical Operation	Sort
Logical Operation	Sort
Actual Execution Mode	Row
Estimated Execution Mode	Row
Actual Number of Rows for All Executions	9
Actual Number of Batches	0
Estimated I/O Cost	0,0112613
Estimated Operator Cost	0,011406 (59%)
Estimated Subtree Cost	0,0146979
Estimated CPU Cost	0,0001448
Estimated Number of Executions	1
Number of Executions	1
Estimated Number of Rows for All Executions	9
Estimated Number of Rows Per Execution	9
Estimated Row Size	261 B
Actual Rebinds	1
Actual Rewinds	0
Node ID	1

**Output List**

[WorkshopDb].[dbo].[ServiceOrders].Id; [WorkshopDb].[dbo].[ServiceOrders].OrderNumber; [WorkshopDb].[dbo].[ServiceOrders].StatusId; [WorkshopDb].[dbo].[ServiceOrders].OpenedAt; [WorkshopDb].[dbo].[ServiceOrders].CustomerNameSnapshot; [WorkshopDb].[dbo].[ServiceOrders].RegistrationNumberSnapshot

**Order By**

[WorkshopDb].[dbo].[ServiceOrders].OpenedAt Descending

### Clustered Index Scan (Clustered)

Scanning a clustered index, entirely or only a range.

Physical Operation	Clustered Index Scan
Logical Operation	Clustered Index Scan
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Actual Number of Rows for All Executions	9
Actual Number of Rows Read	9
Actual Number of Batches	0
Estimated Operator Cost	0,0032919 (17%)
Estimated I/O Cost	0,003125
Estimated CPU Cost	0,0001669
Estimated Subtree Cost	0,0032919
Number of Executions	1
Estimated Number of Executions	1
Estimated Number of Rows to be Read	9
Estimated Number of Rows for All Executions	9
Estimated Number of Rows Per Execution	9
Estimated Row Size	261 B
Actual Rebinds	0
Actual Rewinds	0
Ordered	False
Node ID	2

### Object

[WorkshopDb].[dbo].[ServiceOrders].[PK\_ServiceOrders] [o]

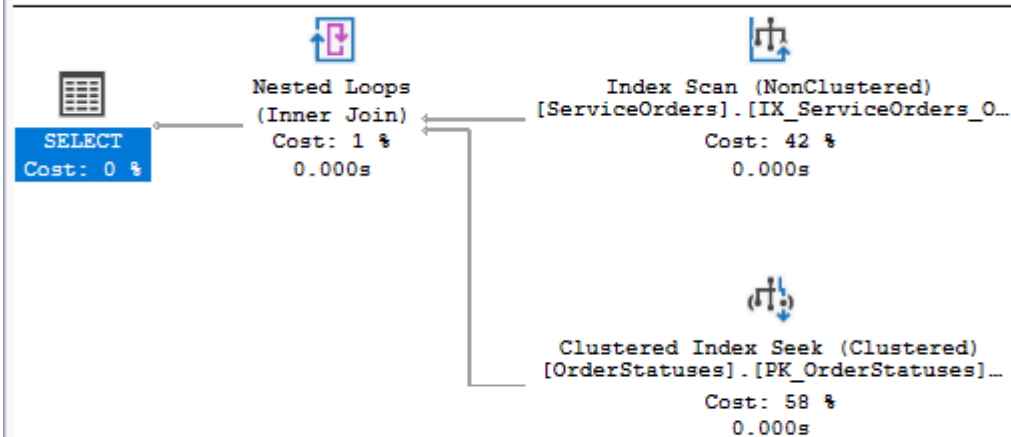
### Output List

[WorkshopDb].[dbo].[ServiceOrders].Id; [WorkshopDb].[dbo].[ServiceOrders].OrderNumber; [WorkshopDb].[dbo].[ServiceOrders].StatusId; [WorkshopDb].[dbo].[ServiceOrders].OpenedAt; [WorkshopDb].[dbo].[ServiceOrders].CustomerNameSnapshot; [WorkshopDb].[dbo].[ServiceOrders].RegistrationNumberSnapshot

Z indeksem na (OpenedAt, StatusId)

Query 1: Query Cost (relative to the batch): 100%

SELECT o.Id, o.OrderNumber, o.OpenedAt, o.CustomerNameSnapshot, o.Reg



Index Scan (NonClustered)	
Scan a nonclustered index, entirely or only a range.	
Physical Operation	Index Scan
Logical Operation	Index Scan
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Actual Number of Rows for All Executions	9
Actual Number of Rows Read	9
Actual Number of Batches	0
Estimated Operator Cost	0,0032919 (42%)
Estimated I/O Cost	0,003125
Estimated CPU Cost	0,0001669
Estimated Subtree Cost	0,0032919
Number of Executions	1
Estimated Number of Executions	1
Estimated Number of Rows to be Read	9
Estimated Number of Rows for All Executions	9
Estimated Number of Rows Per Execution	9
Estimated Row Size	261 B
Actual Rebinds	0
Actual Rewinds	0
Ordered	True
Node ID	1
<b>Object</b>	
[WorkshopDb].[dbo].[ServiceOrders].	
[IX_ServiceOrders_OpenedAt_StatusId] [o]	
<b>Output List</b>	
[WorkshopDb].[dbo].[ServiceOrders].Id; [WorkshopDb].[dbo].	
[ServiceOrders].OrderNumber; [WorkshopDb].[dbo].	
[ServiceOrders].StatusId; [WorkshopDb].[dbo].	
[ServiceOrders].OpenedAt; [WorkshopDb].[dbo].	
[ServiceOrders].CustomerNameSnapshot; [WorkshopDb].[dbo].	
[ServiceOrders].RegistrationNumberSnapshot	

Po dodaniu nieklastrowego indeksu IX\_ServiceOrders\_OpenedAt\_StatusId zapytanie przestało robić pełny skan klastra i „kosztowny sort” – plan zredukował się do nieklastrowego Index Sorta. Natomiast przy naszej ilości danych to i tak praktycznie nic nie dało, gdyż mamy za mało danych.